

COMPUTER ARCHITECTURE

16:332:563

Professor: NAGMEH KARIMI

EVALUATING BRANCH PREDICTION STRATEGIES

Designers

ANEESH ABHYANKAR, ana85

VISHALSINGH HAJERI, vsh15

CONTENTS

| | |
|--|----|
| 1. Abstract..... | 3 |
| 2. Introduction and problem statement..... | 4 |
| 2.1 introduction..... | 4 |
| 2.2 problem statement..... | 4 |
| 3. Software infrastructure..... | 5 |
| 4. Branch prediction strategies..... | 6 |
| 4.1 Static Branch Prediction..... | 6 |
| 4.1.1 Predict Taken..... | 6 |
| 4.1.2 Predict Not Taken..... | 6 |
| 4.2 Dynamic Branch Prediction..... | 6 |
| 4.2.1 1-bit Branch Prediction..... | 7 |
| 4.2.2 2-bit Bimodal Predictor..... | 7 |
| 5. Strategies we implemented..... | 8 |
| 5.1 Strategy 6 – 1-bit Branch Predictor..... | 8 |
| 5.2 Strategy 7 – 2-bit Branch Predictor (Bimodal Predictor)..... | 14 |
| 5.3 Tournament Branch Predictor..... | 20 |
| 5.3.1 Warm up effect..... | 26 |
| 6. Conclusion and Future Scope..... | 28 |
| 6.1 Conclusion..... | 29 |
| 6.2 Future Scope..... | 29 |
| 7. References..... | 30 |

Evaluation of Branch Prediction Strategies

1. ABSTRACT

In order to increase the throughput of a computer we need to increase the number of instructions retiring from it. When the program is sequential fetching next instructions is trivial. But, when there are jumps and branches in the code determining the address of next instructions proves to be a bottleneck in the performance. Exploiting instruction level parallelism becomes crucial in the presence of branch instructions in the form of conditional and unconditional jumps. Branches interrupt the sequential flow of instructions and redirect it to some other location. In this case, fetching sequential instructions results in the eventual flushing of pipeline and re-fetching of instructions. This causes loss of performance in terms of throughput. In superscalar computers, that can fetch multiple instructions in a clock cycle, this loss is more pronounced as all the instructions fetched need to be flushed in the aftermath of a branch instruction. Thus, selecting the best branch prediction strategy is a crucial part of the design of a computer.

2. INTRODUCTION AND PROBLEM STATEMENT

2.1 INTRODUCTION

The ever-increasing need to improve performance has led to extensive research in the field of instruction level parallelism (ILP). The very early solution to the problem of wrongly fetching instructions after branches was to stall the processor until the address of the next instruction was calculated. But this proved to be a serious deterrent to the speed of operation of the processor. Branch prediction strategies provide a way to predict the direction of a branch and fetch the instruction at that location even before the processor calculates the branch target. Static prediction strategies predict a branch direction without considering the dynamic history of the behavior of that branch. Dynamic strategies will consider the recent history of that branch and optionally that of some previous branches before predicting. Sophisticated strategies have increased the processor performance tremendously, by accurately predicting the branch target almost 98% [“Combining Branch Predictors” – June 1993 – Scott McFarling] of the times.

2.2 PROBLEM STATEMENT

In this project we evaluate some famous strategies already published in literature, by implementing them in `simplescalar`. Our problem statement was to evaluate Strategy 6 and Strategy 7 from the paper “Study of Branch Prediction Strategies” by James E. Smith (1981) and a new strategy using `simplescalar` and to juxtapose all of them. The new strategy (not included in `simplescalar`) that we implement is the ‘Tournament Predictor’, which combines local and global predictors. We go ahead and evaluate basic strategies like ‘always taken’, ‘always not taken’, ‘1-bit branch predictor’, ‘2-bit bimodal’, and ‘2 level’ also. We run SPEC95 little endian benchmarks on these strategies and analyze them to find that the new strategy, tournament predictor gives the maximum number of correct predictions, i.e. the minimum rate of misprediction.

3. SOFTWARE INFRASTRUCTURE

We use SimpleScalar toolset, which is a modeling infrastructure that allows users to model and simulate computer microarchitecture for carrying out performance analysis and hardware-software co-verification. The users can, using this tool, simulate various processor environments, complete with static/dynamic scheduling, multiple issues of instructions, non-blocking caches and static/dynamic branch prediction strategies. SimpleScalar can emulate Alpha and Portable Instruction Set Architectures (PISA). We used the PISA, which is a MIPS like architecture, since it is used for instructional and academic purposes. SimpleScalar consists of a cross compiler which basically creates the illusion that the underlying computer is PISA even though it is something else. Todd Austin in the University of Wisconsin at Madison created SimpleScalar, and several other researchers made further contributions to it.

4. BRANCH PREDICTION STRATEGIES

4.1 STATIC BRANCH PREDICTION

They are very primitive kinds of branch prediction strategies, which do not consider the dynamic behavior of the branches during run time of the program. They make a static decision for every branch at compile time, which remains the same during the entire lifetime of the program. They do not have the ability to adapt with the behavior of the program during run time. Examples are taken, not taken branch predictors, etc.

4.1.1 Predict Taken

This strategy always predicts all the branches to be taken and goes ahead with this prediction to fetch the next instruction at the branch target. This technique is found to work better with loop instructions, since they are almost always taken so as to loop back to the next iteration. It is only once after completion of the loop that the branch will be not taken. The problem with this strategy though is that the processor needs to calculate the branch target before checking the branch condition. This requires the branch target calculation hardware, i.e. 'adder' to be placed in the fetch cycle of the execution pipeline.

4.1.2 Predict Not Taken

This strategy will predict all branches to be not taken and carry the execution to the sequentially next instruction. This strategy fails miserably with loops for the reason mentioned above. Normally a typical program is made up of loops and this makes this strategy a bad choice in normal circumstances.

4.2 DYNAMIC BRANCH PREDICTION

These strategies use the information collected about a branch at run time to make a prediction about its direction. Unlike static predictors, which are compiler based, dynamic predictors are hardware implementations and perform a lot better than static predictors. The reason is mainly the fact that they take into consideration the dynamic behavior of branches while predicting their direction, which helps them correct a wrong decision taken in the past. The predictor may also become more inclined to make a particular prediction for a particular branch if that branch produces consistent outcomes. Examples are 1-bit branch predictor, 2-bit bimodal predictors, etc.

4.2.1 1-bit Branch Predictor

This is a dynamic prediction strategy, which maintains a 1-bit saturating counter. It has a history table that is indexed into by lower few bits of the Program Counter (PC). Entries of this table are 1-bit counters that help make predictions.

4.2.2 2-bit Bimodal Predictor

The problem with 1-bit Branch Predictor is that it changes its prediction very quickly. When a particular branch direction is mispredicted, the strategy changes its prediction for the next occurrence of that branch. Such a predictor fails when a branch is alternately taken. For example, for a program that prints all even numbers, the branch is taken alternately and this results in a very low accuracy. A simple solution to this problem is to provide hysteresis to the 1-bit predictor and to not change the prediction very quickly. 2-bit bimodal predictor uses a 2-bit saturating counter instead of a 1-bit counter in the earlier strategy. When the counter is 0 or 1, the branch is predicted not taken and when the counter is 2 or 3, the branch is predicted taken.

5. STRATEGIES WE IMPLEMENTED

5.1 STRATEGY 6 – 1-BIT BRANCH PREDICTOR

Strategy 6 described in the paper is like 1-bit branch predictor explained above. This strategy is implemented in simplescalar by defining a 1-bit saturating counter and updating it as per the figure 1. When the value of the counter is 1, the branch is predicted taken and if this proves to be a misprediction, the counter is decremented. When the counter value is 0, the branch is predicted not taken and the counter is incremented in the event of misprediction. As shown in Figure 2, the lower ‘M’ bits of PC are used to index into the prediction table that contains such 1-bit saturating counters in each entry.




Figure 1 – State Machine of 1-bit Branch Predictor (courtesy: www.umd.edu)




Figure 2 – 1-bit Branch Predictor

We increase the prediction table size in the predictor and obtain the number of prediction hits to study the effect of table size on the prediction accuracy. We see as expected that the prediction accuracy goes on increasing as we increase table size. Increasing table size implicitly means increasing the number of lower order bits of PC used to hash into the table. Results of the analysis are illustrated in figures 3 to 7 below.

| 1 bit Branch Predictor - Compress95 | | |
|-------------------------------------|----------|------------|
| Number of Branches | 10992983 | |
| | | |
| Table Size | Misses | Hits |
| 8 | 3440414 | 7552569 |
| 16 | 2763397 | 8229586 |
| 32 | 2351026 | 8641957 |
| 64 | 2350273 | 8642710 |
| 128 | 2347043 | 8645940 |
| 256 | 2345924 | 8647059 |
| 512 | 2344839 | 8648144 |
| 1024 | 2344464 | 8648519 |
| 2048 | 2344366 | 8648617 |
| 4096 | 2343704 | 8649279 |
| | | |
| Maximum Hits | | 8649279 |
| Accuracy | | 78.6799998 |




Figure 3 – Analysis of Strategy 6 against Compress95 Benchmark

| 1 bit Branch Predictor - GO | | |
|-----------------------------|----------|----------|
| Number of Branches | 71660927 | |
| | | |
| Table Size | Misses | Hits |
| 8 | 25981954 | 45678973 |
| 16 | 26327755 | 45333172 |
| 32 | 25986817 | 45674110 |
| 64 | 24231012 | 47429915 |
| 128 | 23111114 | 48549813 |
| 256 | 22179887 | 49481040 |
| 512 | 21391393 | 50269534 |
| 1024 | 20729531 | 50931396 |
| 2048 | 20400215 | 51260712 |
| 4096 | 20100890 | 51560037 |
| | | |
| Maximum Hits | | 51560037 |
| Accuracy | | 71.95 |

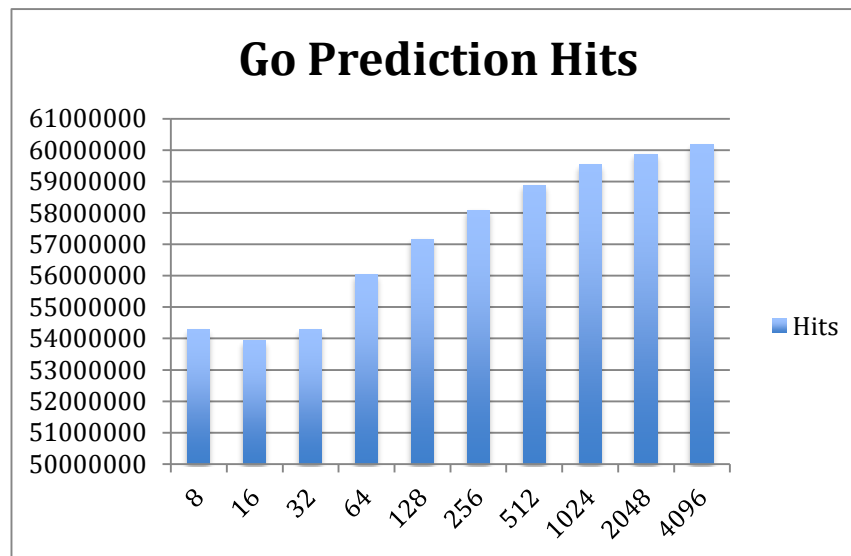


Figure 4 – Analysis of Strategy 6 against Go Benchmark

| 1 bit Branch Predictor - PERL | | |
|-------------------------------|---------|------------|
| Number of Branches | 3061732 | |
| | | |
| Table Size | Misses | Hits |
| 8 | 1374018 | 1687714 |
| 16 | 1244166 | 1817566 |
| 32 | 1052040 | 2009692 |
| 64 | 972603 | 2089129 |
| 128 | 841651 | 2220081 |
| 256 | 732944 | 2328788 |
| 512 | 614723 | 2447009 |
| 1024 | 556451 | 2505281 |
| 2048 | 529555 | 2532177 |
| 4096 | 512840 | 2548892 |
| | | |
| Maximum Hits | | 2548892 |
| Accuracy | | 83.2500036 |




Figure 5 – Analysis of Strategy 6 against Perl Benchmark

| 1 bit Branch Predictor - CC1 | | |
|------------------------------|----------|------------|
| Number of Branches | 41277100 | |
| | | |
| Local History Size | Misses | Hits |
| 8 | 17492356 | 23784744 |
| 16 | 16420548 | 24856552 |
| 32 | 15189544 | 26087556 |
| 64 | 13785887 | 27491213 |
| 128 | 12137649 | 29139451 |
| 256 | 10752545 | 30524555 |
| 512 | 9614366 | 31662734 |
| 1024 | 8758463 | 32518637 |
| 2048 | 8140064 | 33137036 |
| 4096 | 7727072 | 33550028 |
| | | |
| Maximum Hits | | 33550028 |
| Accuracy | | 81.2800027 |




Figure 6 – Analysis of Strategy 6 against CC1 Benchmark

| 1 bit Branch Predictor – ANAGRAM | | |
|----------------------------------|--------|------------|
| Number of Branches | 757 | |
| | | |
| Table Size | Misses | Hits |
| 8 | 197 | 560 |
| 16 | 187 | 570 |
| 32 | 182 | 575 |
| 64 | 180 | 577 |
| 128 | 177 | 580 |
| 256 | 172 | 585 |
| 512 | 176 | 581 |
| 1024 | 172 | 585 |
| 2048 | 172 | 585 |
| 4096 | 174 | 583 |
| | | |
| Maximum Hits | | 585 |
| Accuracy | | 77.2787318 |




Figure 7 – Analysis of Strategy 6 against Anagram Benchmark

5.2 STRATEGY 7 – 2-BIT BRANCH PREDICTOR (BIMODAL PREDICTOR)

As discussed in section 4.2.2, 1-bit branch predictor faces problems when the actual branch outcomes change from taken to not taken frequently. This strategy provides an improvement over the 1-bit strategy by implementing 2-bit counters to take prediction decisions. The state machine implemented by this strategy to update the counter values is shown in figure 3 and the branch predictor algorithm is shown in figure 4.

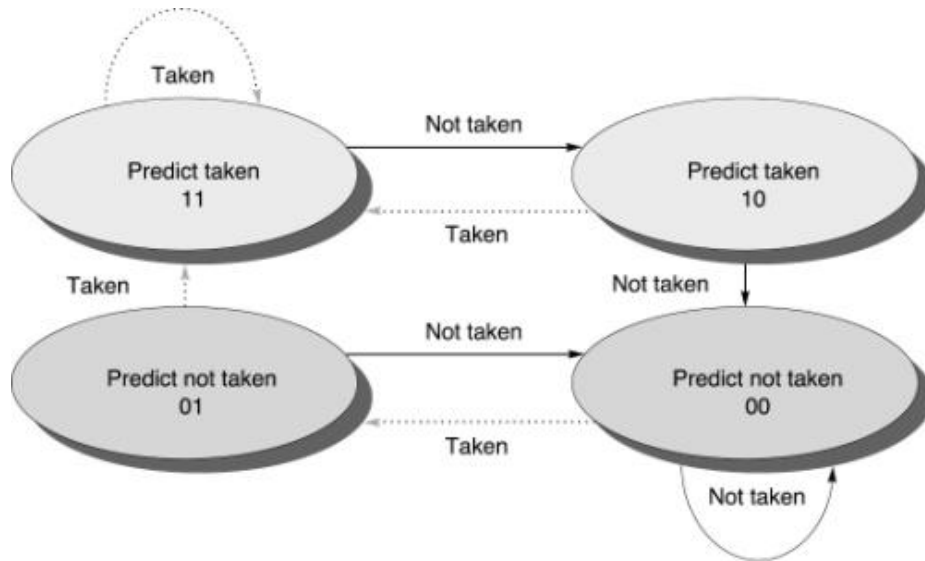


Figure 8 – State Machine of 2-bit Branch Predictor (courtesy: www.ncsu.edu)

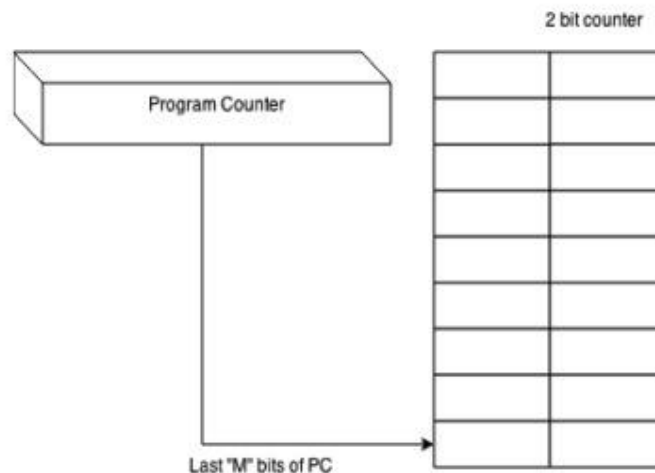


Figure 9 – 2-bit Branch Predictor (Bimodal Predictor)

We carry out the same analysis for 2-bit branch predictor as we did for 1-bit branch predictor to find the accuracy. We see that the performance of 2-bit predictor is better than that of 1-bit predictor as it adds hysteresis to it. The analysis has been illustrated in figures 10 to 14 below.

| 2 bit Branch Predictor - Compress95 | | |
|-------------------------------------|----------|------------|
| Number of Branches | 14361040 | |
| | | |
| Table Size | Misses | Hits |
| 8 | 1887445 | 11989635 |
| 16 | 1607468 | 12473595 |
| 32 | 1454012 | 12753572 |
| 64 | 1450936 | 12907028 |
| 128 | 1449448 | 12910104 |
| 256 | 1448817 | 12911592 |
| 512 | 1448274 | 12912223 |
| 1024 | 1448067 | 12912766 |
| 2048 | 1447658 | 12912973 |
| 4096 | 2343704 | 12913382 |
| | | |
| Maximum Hits | | 12913382 |
| Accuracy | | 89.9195462 |

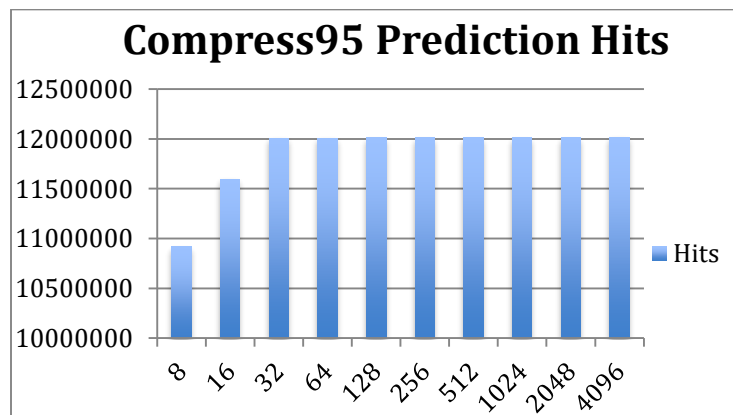


Figure 10 – Analysis of Strategy 7 against Compress95 Benchmark

| 2 bit Branch Predictor - GO | | |
|-----------------------------|----------|------------|
| Number of Branches | 80274754 | |
| | | |
| Table Size | Misses | Hits |
| 8 | 24869141 | 55405613 |
| 16 | 24816805 | 55457949 |
| 32 | 23832433 | 56442321 |
| 64 | 22355499 | 57919255 |
| 128 | 20640618 | 59634136 |
| 256 | 18919622 | 61355132 |
| 512 | 17630318 | 62644436 |
| 1024 | 16741821 | 63532933 |
| 2048 | 16039005 | 64235749 |
| 4096 | 15515991 | 64758763 |
| | | |
| Maximum Hits | | 64758763 |
| Accuracy | | 80.6713939 |




Figure 11 – Analysis of Strategy 7 against GO Benchmark

| 2 bit Branch Predictor - PERL | | |
|-------------------------------|---------|-----------|
| Number of Branches | 4366187 | |
| | | |
| Table Size | Misses | Hits |
| 8 | 1264387 | 3101800 |
| 16 | 1139814 | 3226373 |
| 32 | 972756 | 3393431 |
| 64 | 818882 | 3547305 |
| 128 | 701926 | 3664261 |
| 256 | 615171 | 3751016 |
| 512 | 474220 | 3891967 |
| 1024 | 421029 | 3945158 |
| 2048 | 388377 | 3977810 |
| 4096 | 370076 | 3996111 |
| | | |
| Maximum Hits | | 3996111 |
| Accuracy | | 91.524046 |




Figure 12 – Analysis of Strategy 7 against PERL Benchmark

| 2 bit Branch Predictor - Compress95 | | |
|-------------------------------------|----------|------------|
| Number of Branches | 56338415 | |
| | | |
| Table Size | Misses | Hits |
| 8 | 16036651 | 40301764 |
| 16 | 14794429 | 41543986 |
| 32 | 13507026 | 42831389 |
| 64 | 11985020 | 44353395 |
| 128 | 10378785 | 45959630 |
| 256 | 8868439 | 47469976 |
| 512 | 7748289 | 48590126 |
| 1024 | 6917451 | 49420964 |
| 2048 | 6317122 | 50021293 |
| 4096 | 5934658 | 50403757 |
| | | |
| Maximum Hits | | 50403757 |
| Accuracy | | 89.4660544 |




Figure 13 – Analysis of Strategy 7 against CC1 Benchmark

| 2 bit Branch Predictor - Compress95 | | |
|-------------------------------------|--------|------------|
| Number of Branches | 971 | |
| | | |
| Table Size | Misses | Hits |
| 8 | 186 | 785 |
| 16 | 178 | 793 |
| 32 | 164 | 807 |
| 64 | 165 | 806 |
| 128 | 146 | 825 |
| 256 | 152 | 819 |
| 512 | 147 | 824 |
| 1024 | 144 | 827 |
| 2048 | 145 | 826 |
| 4096 | 144 | 827 |
| | | |
| Maximum Hits | | 827 |
| Accuracy | | 89.9195462 |




Figure 14 – Analysis of Strategy 7 against Anagram Benchmark

5.3 TOURNAMENT BRANCH PREDICTOR

The predictors analyzed above are local predictors, which means that they only consider the local branch history. They only consider the history of the branch of which direction is to be predicted. But, many times the outcome of a branch is not only dependent on the history of that branch, but also on the history of previous branches. This observation gave rise to the researchers developing global branch predictors, which provide consideration of the same. These predictors intuitively perform better than local predictors. However, neither global nor local predictors are good enough when we require accuracies in the range of 98%.

So, they developed a new type of a predictor called tournament predictor. The idea that is to introduce a predictor to predict which predictor will predict better. It combines the best performance of both local and global predictor. It consists of a local predictor table, a global predictor table and a selector to select between the two. So, the selector is a 2-bit saturating counter, which has the following states: 00 – strongly prefer local predictor, 01 – weakly prefer local predictor, 10 – weakly prefer global predictor and 11 – strongly prefer global predictor. The tournament predictor implemented with configuration: <L1size = 1024, L2 size = 1024, GHR size = 8, XOR = Yes> is shown in figure 15 below.

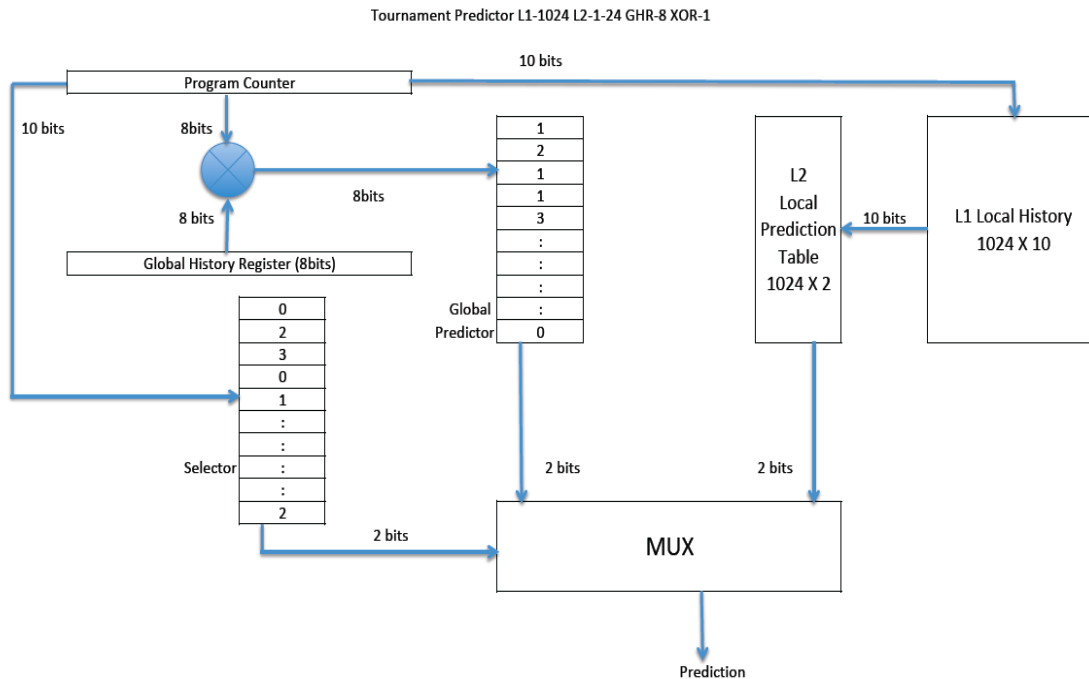


Figure 15 - Tournament Branch Predictor

| Tournament Predictor - Compress95 | | |
|-----------------------------------|----------|-----------------|
| Number of Branches | 14361040 | |
| | | |
| Config: (variable) 1024 8 1 | Misses | Compress95-Hits |
| 1024 | 1033548 | 13326879 |
| 512 | 1034161 | 13326879 |
| 256 | 1034243 | 13326797 |
| 128 | 1035416 | 13325624 |
| 64 | 1037055 | 13323985 |
| | | |
| Config: 1024 1024 (variable) 1 | Misses | Compress95-Hits |
| 4 | 1039536 | 13321504 |
| 5 | 896729 | 13464311 |
| 6 | 872487 | 13488553 |
| 7 | 936863 | 13424177 |
| 8 | 1033548 | 13327492 |
| 9 | 1090909 | 13270131 |
| 10 | 1099582 | 13261458 |

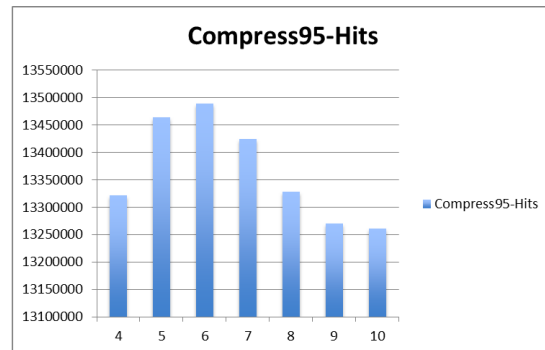



Figure 16 – Analysis of Tournament predictor against Compress95 Benchmark

| Tournament Predictor - GO | | |
|-------------------------------|----------|----------|
| Number of Branches | 80274754 | |
| | | |
| Config (variable) 1024 8 1 | Misses | GO-Hits |
| 4096 | 17948546 | 62326208 |
| 2048 | 18402108 | 61872646 |
| 1024 | 18977828 | 61296926 |
| 512 | 19663010 | 60611744 |
| 256 | 20622699 | 59652055 |
| 128 | 21584673 | 58690081 |
| 64 | 21998642 | 58276112 |
| | | |
| Config 4096 1024 (variable) 1 | Misses | GO-Hits |
| 4 | 17538823 | 62735931 |
| 5 | 17354122 | 62920632 |
| 6 | 15141708 | 65133046 |
| 7 | 16542892 | 63731862 |
| 8 | 17948546 | 62326208 |
| 9 | 18267429 | 62007325 |
| 10 | 18444566 | 61830188 |

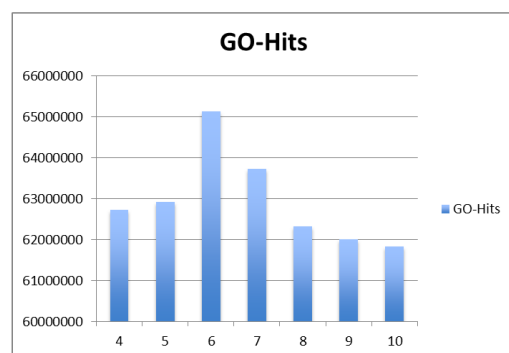
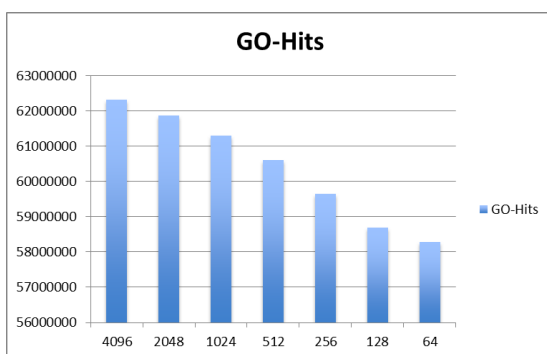




Figure 17 – Analysis of Tournament predictor against GO Benchmark

| Tournament Predictor - PERL | | |
|-------------------------------|---------|-----------|
| Number of Branches | 4366187 | |
| | | |
| Config (variable) 1024 8 1 | Misses | PERL-Hits |
| 4096 | 468195 | 3897992 |
| 2048 | 483214 | 3882973 |
| 1024 | 500842 | 3865345 |
| 512 | 542404 | 3823783 |
| 256 | 645678 | 3720509 |
| 128 | 765438 | 3600749 |
| 64 | 816504 | 3549683 |
| | | |
| Config 4096 1024 (variable) 1 | Misses | PERL-Hits |
| 4 | 459853 | 3906334 |
| 5 | 437131 | 3929056 |
| 6 | 362770 | 4003417 |
| 7 | 404893 | 3961294 |
| 8 | 468195 | 3897992 |
| 9 | 519853 | 3846334 |
| 10 | 525583 | 3840604 |



18 – Analysis of Tournament predictor against Perl Benchmark

| Tournament Predictor - CC1 | | |
|-------------------------------|----------|----------|
| Number of Branches | 56338415 | |
| | | |
| Config (variable) 1024 8 1 | Misses | CC1-Hits |
| 4096 | 7030925 | 49307490 |
| 2048 | 7289391 | 49049024 |
| 1024 | 7730431 | 48607984 |
| 512 | 8353360 | 47985055 |
| 256 | 9356793 | 46981622 |
| 128 | 10875931 | 45462484 |
| 64 | 11500972 | 44837443 |
| | | |
| Config 4096 1024 (variable) 1 | Misses | CC1-Hits |
| 4 | 6948023 | 49390392 |
| 5 | 6501195 | 49837220 |
| 6 | 5769734 | 50568681 |
| 7 | 6208653 | 50129762 |
| 8 | 7030925 | 49307490 |
| 9 | 7156983 | 49181432 |
| 10 | 7293613 | 49044802 |


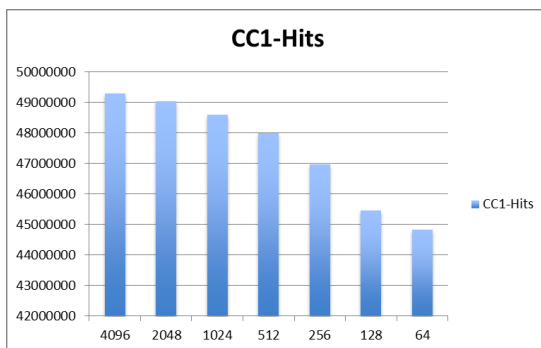


Figure 19 – Analysis of Tournament predictor against CC1 Benchmark

| Tournament Predictor - Anagram | | |
|--------------------------------|--------|--------------|
| | 971 | |
| Config = 1024 1024 8 1 | | |
| | | |
| Config (variable) 1024 8 1 | Misses | Anagram-Hits |
| 1024 | 159 | 812 |
| 512 | 170 | 801 |
| 256 | 177 | 794 |
| 128 | 185 | 786 |
| 64 | 193 | 778 |
| | | |
| Config 1024 1024 (variable) 1 | Misses | Anagram-Hits |
| 4 | 171 | 800 |
| 5 | 166 | 805 |
| 6 | 159 | 812 |
| 7 | 160 | 811 |
| 8 | 155 | 816 |
| 9 | 169 | 802 |
| 10 | 176 | 795 |





Figure 20 – Analysis of Tournament predictor against Anagram Benchmark

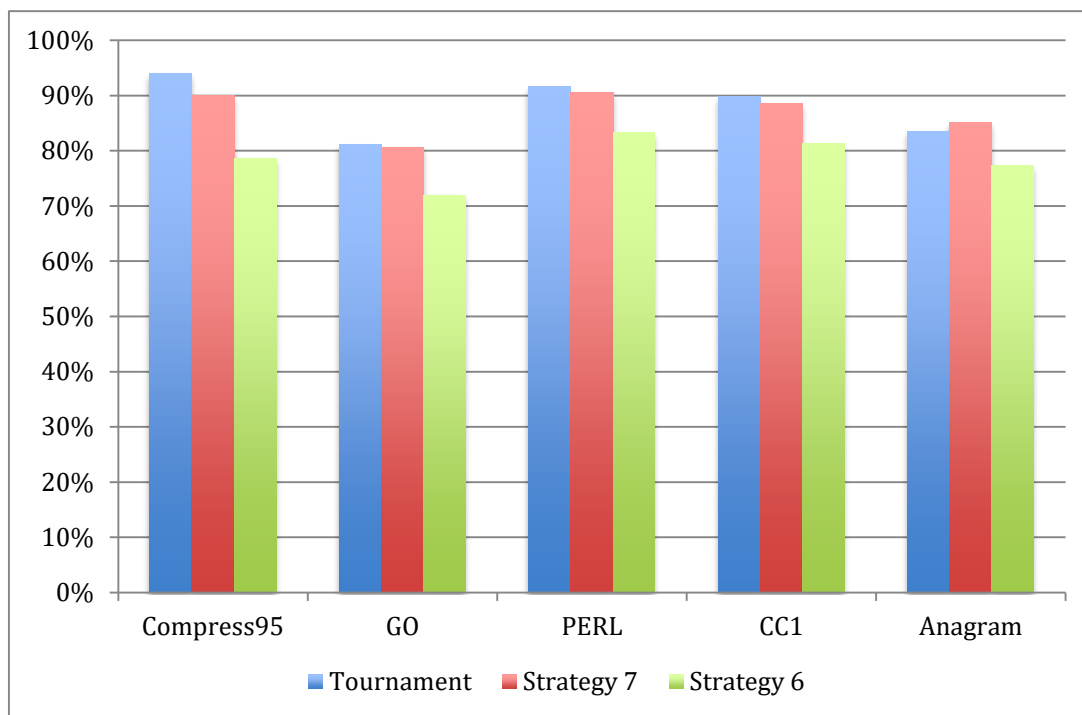


Figure 21 – Comparison of the performance of all Branch predictors on all Benchmarks

5.3.1 Warm-up Effect

As seen from the results of three different strategies, tournament predictor performs slightly less than strategy 7. But, as per our understanding since tournament predictor is a combination of local and global predictor, while 2-bit bimodal is just one level local predictor, tournament predictor should always perform better than strategy 7.

We dug deeper to understand this anomalous behavior of branch prediction strategies in case of Anagram benchmark. Anagram has only 971 branch instructions. The shift registers used in local predictor and branch history register, which store outcomes of previous predictions are initialized to some random bits (all 0's or all 1's) at the start of the program. So, Tournament predictor mispredicts initial few branch instructions as its predictions are not based upon any run-time data.

This warm-up effect is more evident in case of tournament predictors than Bimodal predictors. As seen in the graph, accuracy of tournament predictor is lower for fewer number of instructions, but as the total number of branch instructions goes on increasing

so does the accuracy of tournament predictor. In the end, tournament predictor settles down to a much higher accuracy as compared to Bimodal predictor.




Figure 22 – Visualizing the effect of Warm up

6. CONCLUSION AND FUTURE SCOPE

6.1 Conclusion

As it is clearly seen from Figure 21 that Strategy 7 which uses 2-bit saturating counter performs better than strategy 6 which uses 1-bit saturating counter. The increased accuracy in strategy 7 can be contributed to one of its important feature of not being as impulsive as strategy 6 in making decisions.

Conditional branches encountered in the program could be either correlating branches i.e whose direction depend upon the outcomes of some 'n' previous branches or it could be non-correlating. Our understanding of branch predictor literature tells us that local branch predictors work well in case of non-correlating branches, while global predictors work well with correlating branches. Thus, our implemented strategy produces better results in comparison to previous strategies by combining both local and global branch predictor. It chooses one among the two predictors depending on the type of branch.

The one shortcoming of tournament predictor is that its accuracy is low for a benchmark having less number of branch instructions (<10000) because of warm-up effect. The w-bit shift registers used in local predictor and branch history register is initialized to some random bits (all 0's or all 1's). So, Tournament predictor mispredicts initial few branch instructions as it doesn't have run-time data to make predictions upon.

In tournament predictor, the choice of global predictor is crucial. We chose gshare global predictor over gselect global predictor since gshare XOR's branch PC and global branch history registers while gselect concatenates the same. XORing minimizes the effects of destructive aliasing, which are evident in gselect.

The Tournament predictor performs better than any single level branch predictor provided that there are higher number of branch instructions in the code.

6.2 Future Scope

An alternative to branch prediction could be executing both the directions of a branch speculatively. However, required hardware resources will increase and will be proportional to number of concurrent speculative executions.

7. REFERENCES

- A Study of Branch Prediction Strategies – James Smith
- Combining Branch Predictors – June 1993 – Scott McFarling
- Simple scalar user guide: <http://www.simplescalar.com/>
- http://ocw.mit.edu/courses/electrical-engineering-and-computer-science/6-823-computer-system-architecture-fall-2005/lecture-notes/113_brnchpred.pdf
- <http://users.ece.gatech.edu/~sudha/academic/class/ece4100->
- Smith JE. A study of branch prediction strategies. Proceedings of the 8th Annual International Symposium on Computer Architecture, Minneapolis, MN, 12–14 May 1981. ACM Press, 1981; 135–148.
- Hennessy J, Patterson D. Computer Architecture: A Quantitative Approach. Morgan Kaufmann: San Mateo, CA, 2003
- Nair R. Dynamic path-based branch correlation. Proceedings of the 28th Annual International Symposium on Microarchitecture, Ann Arbor, MI, 29 November–1 December 1995. ACM Press, 1995; 15–23.